S2 Figure: Domain-Stratified Benchmark Results

Candidate number: 1102384

Figure S2 displays paired differences separately for all seven benchmark task groups.

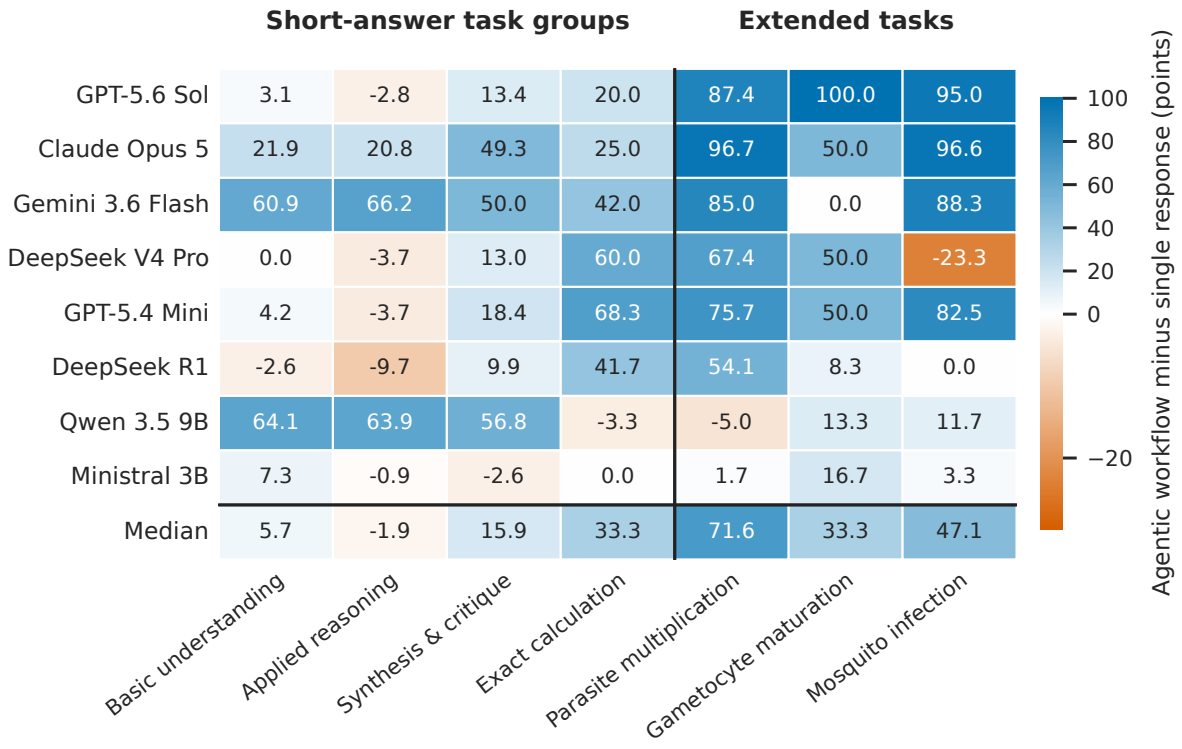

Figure S2: Paired score differences by benchmark task group. Each cell gives the agentic-workflow minus single-response mean for one model and task group; the final row gives descriptive medians across the eight fixed models. Positive values favour the agentic workflow. Short-answer differences are percentage points and extended-task differences are benchmark points.